

```
In [24]: import pandas as pd
import numpy as np
import glob
import xgboost as xgb
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.preprocessing import LabelEncoder
from xgboost import XGBClassifier
from sklearn.utils.class_weight import compute_sample_weight
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import precision_recall_fscore_support
```

Data is present in Epl/data/ and separate features required are taken to reduce the difficulties in cleaning..empty columns if any are filled with NaN.

```
In [25]: # Load only Premier League CSVs (E0)
files = glob.glob("C:/Epl/data/E0*.csv")
print(files)
features=["Date", "HomeTeam", "AwayTeam", "HTHG", "HTAG", "HS", "AS", "HST", "AST", "HF", "AF", "HC", "AC", "HY", "AY", "HR", "AR", "FTR"]
target = "FTR"
dfs = []
for f in files:
    df = pd.read_csv(f, encoding='cp1252', on_bad_lines='skip')
    for col in features:
        if col not in df.columns:
            df[col] = pd.NA
    df = df[features + [target]]
    dfs.append(df)

data = pd.concat(dfs, ignore_index=True)
data["Date"] = pd.to_datetime(data["Date"], dayfirst=True, errors="coerce")
print("Shape:", data.shape)
print(data.head())
```

```
['C:/Epl/data\\E02000.csv', 'C:/Epl/data\\E02001.csv', 'C:/Epl/data\\E02002.csv', 'C:/Epl/data\\E02003.csv', 'C:/Epl/data\\E02004.csv', 'C:/Epl/data\\E02005.csv', 'C:/Epl/data\\E02006.csv', 'C:/Epl/data\\E02007.csv', 'C:/Epl/data\\E02008.csv', 'C:/Epl/data\\E02009.csv', 'C:/Epl/data\\E02010.csv', 'C:/Epl/data\\E02011.csv', 'C:/Epl/data\\E02012.csv', 'C:/Epl/data\\E02013.csv', 'C:/Epl/data\\E02014.csv', 'C:/Epl/data\\E02015.csv', 'C:/Epl/data\\E02016.csv', 'C:/Epl/data\\E02017.csv', 'C:/Epl/data\\E02018.csv', 'C:/Epl/data\\E02019.csv', 'C:/Epl/data\\E02020.csv', 'C:/Epl/data\\E02021.csv', 'C:/Epl/data\\E02022.csv', 'C:/Epl/data\\E02023.csv', 'C:/Epl/data\\E02024.csv']
```

Shape: (9411, 18)

	Date	HomeTeam	AwayTeam	HTHG	HTAG	HS	AS	HST	AST	\
0	2000-08-19	Charlton	Man City	2.0	0.0	17.0	8.0	14.0	4.0	
1	2000-08-19	Chelsea	West Ham	1.0	0.0	17.0	12.0	10.0	5.0	
2	2000-08-19	Coventry	Middlesbrough	1.0	1.0	6.0	16.0	3.0	9.0	
3	2000-08-19	Derby	Southampton	1.0	2.0	6.0	13.0	4.0	6.0	
4	2000-08-19	Leeds	Everton	2.0	0.0	17.0	12.0	8.0	6.0	

	HF	AF	HC	AC	HY	AY	HR	AR	FTR
0	13.0	12.0	6.0	6.0	1.0	2.0	0.0	0.0	H
1	19.0	14.0	7.0	7.0	1.0	2.0	0.0	0.0	H
2	15.0	21.0	8.0	4.0	5.0	3.0	1.0	0.0	A
3	11.0	13.0	5.0	8.0	1.0	1.0	0.0	0.0	D
4	21.0	20.0	6.0	4.0	1.0	3.0	0.0	0.0	H

C:\Users\arnas\AppData\Local\Temp\ipykernel_20292\2351140533.py:16: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
data["Date"] = pd.to_datetime(data["Date"], dayfirst=True, errors="coerce")
```

Now that data is ready we have to clean it by checking for duplicates and null values.

```
In [26]: print(data.isnull().sum())
print((data.isnull().mean()*100).round(2))
```

```

Date          1
HomeTeam      1
AwayTeam      1
HTHG          1
HTAG          1
HS            1
AS            1
HST           1
AST           1
HF            1
AF            1
HC            1
AC            1
HY            1
AY            1
HR            1
AR            1
FTR           1
dtype: int64
Date          0.01
HomeTeam      0.01
AwayTeam      0.01
HTHG          0.01
HTAG          0.01
HS            0.01
AS            0.01
HST           0.01
AST           0.01
HF            0.01
AF            0.01
HC            0.01
AC            0.01
HY            0.01
AY            0.01
HR            0.01
AR            0.01
FTR           0.01
dtype: float64

```

This shows that the dataset has only a very few negligible output there is 1 missing value and that missing value exists in all features so we could drop that row.

```

In [27]: data = data.dropna()
print("After dropping missing:",data.shape)
print("Remaining NaN:",data.isnull().sum().sum())

```

```

After dropping missing: (9410, 18)
Remaining NaN: 0

```

```

In [28]: duplicates = data.duplicated().sum()
print("Number of duplicated rows:",duplicates)

```

```

Number of duplicated rows: 0

```

```

In [29]: print(data.columns)
le = LabelEncoder()
data['FTR']=le.fit_transform(data['FTR'])
data['Date'] = data['Date'].astype(str)

```

```

Index(['Date', 'HomeTeam', 'AwayTeam', 'HTHG', 'HTAG', 'HS', 'AS', 'HST',
       'AST', 'HF', 'AF', 'HC', 'AC', 'HY', 'AY', 'HR', 'AR', 'FTR'],
      dtype='object')

```

I have done label encoding for all categorical columns and the encoder function is also saved in encoders[] to decode for further use

```

In [30]: import pandas as pd
import numpy as np

data['HT_goal_diff'] = data['HTHG'] - data['HTAG']

data['corner_diff'] = data['HC'] - data['AC']

data['home_shot_accuracy'] = data['HST'] / data['HS'].replace(0, 1)
data['away_shot_accuracy'] = data['AST'] / data['AS'].replace(0, 1)

data['cards_diff'] = ((data['HY'] + 2*data['HR']) -
                      (data['AY'] + 2*data['AR']))

data['fouls_diff'] = data['HF'] - data['AF']

```

```
data['shots_diff'] = data['HS'] - data['AS']

data['shots_on_target_diff'] = data['HST'] - data['AST']
```

In [31]: *#splitting data into label and feature input*

```
y = data['FTR']

X = data.drop('FTR', axis=1)
```

In [32]: `print(data.head())`

```
      Date HomeTeam AwayTeam HTHG HTAG HS AS HST AST \
0  2000-08-19  Charlton   Man City    2.0  0.0  17.0  8.0  14.0  4.0
1  2000-08-19  Chelsea   West Ham    1.0  0.0  17.0  12.0  10.0  5.0
2  2000-08-19  Coventry  Middlesbrough  1.0  1.0  6.0  16.0  3.0  9.0
3  2000-08-19   Derby   Southampton  1.0  2.0  6.0  13.0  4.0  6.0
4  2000-08-19   Leeds   Everton    2.0  0.0  17.0  12.0  8.0  6.0

      HF ... AR FTR HT_goal_diff corner_diff home_shot_accuracy \
0  13.0 ... 0.0  2      2.0          0.0          0.823529
1  19.0 ... 0.0  2      1.0          0.0          0.588235
2  15.0 ... 0.0  0      0.0          4.0          0.500000
3  11.0 ... 0.0  1     -1.0         -3.0          0.666667
4  21.0 ... 0.0  2      2.0          2.0          0.470588

      away_shot_accuracy cards_diff fouls_diff shots_diff \
0          0.500000      -1.0          1.0          9.0
1          0.416667      -1.0          5.0          5.0
2          0.562500          4.0         -6.0         -10.0
3          0.461538          0.0         -2.0          -7.0
4          0.500000      -2.0          1.0          5.0

      shots_on_target_diff
0          10.0
1           5.0
2          -6.0
3          -2.0
4           2.0
```

[5 rows x 26 columns]

Splitting Data

In [45]:

```
from sklearn.model_selection import train_test_split
# Stratified split to preserve class ratios
X_train_cat_split, X_test_cat_split, y_train_cat_split, y_test_cat_split = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

CATBOOST CLASSIFIER

Fitting with catboost

In [46]:

```
# Define which columns are categorical
from catboost import CatBoostClassifier
cat_features = ['HomeTeam', 'AwayTeam', 'Date']

# Build model
model = CatBoostClassifier(
    iterations=1000,
    learning_rate=0.05,
    depth=7,
    loss_function='MultiClass',
    auto_class_weights='Balanced',
    early_stopping_rounds=50,
    verbose=100
)

# Train model
model.fit(X_train_cat_split, y_train_cat_split, cat_features=cat_features, eval_set=(X_test_cat_split, y_test_cat_split))
```

```

0:      learn: 1.0747858      test: 1.0755246 best: 1.0755246 (0)      total: 41.9ms      remaining: 41.8s
100:    learn: 0.7536707      test: 0.8081766 best: 0.8081442 (99)     total: 4.81s      remaining: 42.8s
200:    learn: 0.7035799      test: 0.8013977 best: 0.8011503 (163)    total: 9.44s      remaining: 37.5s
300:    learn: 0.6561751      test: 0.8014106 best: 0.8001670 (277)    total: 13.9s      remaining: 32.2s
Stopped by overfitting detector (50 iterations wait)

```

```

bestTest = 0.8001669608
bestIteration = 277

```

Shrink model to first 278 iterations.

```
Out[46]: <catboost.core.CatBoostClassifier at 0x27d7b758050>
```

Predicting and accuracy

```

In [47]: y_pred = model.predict(X_test_cat_split)

# Accuracy
print("Accuracy:", accuracy_score(y_test_cat_split, y_pred))

# Full classification report
print(classification_report(y_test_cat_split, y_pred))

```

```

Accuracy: 0.6301806588735388
      precision    recall  f1-score   support

     0:       0.68       0.69       0.69         556
     1:       0.37       0.46       0.41         464
     2:       0.79       0.68       0.73         862

 accuracy          0.63         1882
 macro avg       0.61       0.61       0.61         1882
 weighted avg    0.65       0.63       0.64         1882


```

```

In [58]: cm = confusion_matrix(y_test_cat_split, y_pred)

# Plot confusion matrix
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Class 0', 'Class 1', 'Class 2'],
            yticklabels=['Class 0', 'Class 1', 'Class 2'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

```



```


In [59]: prec, rec, f1, _ = precision_recall_fscore_support(y_test_cat_split, y_pred, average=None)

metrics_df = {
    "Class": ['0', '1', '2'],
    "Precision": prec,
    "Recall": rec,
    "F1-Score": f1
}

```

```
metrics_df = pd.DataFrame(metrics_df)
```

```
metrics_df.plot(x="Class", kind="bar", figsize=(8,6))
plt.title("Per-Class Metrics")
plt.ylabel("Score")
plt.xticks(rotation=0)
plt.ylim(0,1)
plt.show()
```



```
In [37]: data = data.drop(columns=["Date"])
categorical_col = data.select_dtypes(include='object').columns
encoders = {}
for c in categorical_col:
    le = LabelEncoder()
    data[c]=le.fit_transform(data[c])
    encoders[c]=le
print(data.head())
```

	HomeTeam	AwayTeam	HTHG	HTAG	HS	AS	HST	AST	HF	AF	...	\
0	12	26	2.0	0.0	17.0	8.0	14.0	4.0	13.0	12.0	...	
1	13	43	1.0	0.0	17.0	12.0	10.0	5.0	19.0	14.0	...	
2	14	28	1.0	1.0	6.0	16.0	3.0	9.0	15.0	21.0	...	
3	16	36	1.0	2.0	6.0	13.0	4.0	6.0	11.0	13.0	...	
4	22	17	2.0	0.0	17.0	12.0	8.0	6.0	21.0	20.0	...	

	AR	FTR	HT_goal_diff	corner_diff	home_shot_accuracy	\
0	0.0	2	2.0	0.0	0.823529	
1	0.0	2	1.0	0.0	0.588235	
2	0.0	0	0.0	4.0	0.500000	
3	0.0	1	-1.0	-3.0	0.666667	
4	0.0	2	2.0	2.0	0.470588	

	away_shot_accuracy	cards_diff	fouls_diff	shots_diff	\
0	0.500000	-1.0	1.0	9.0	
1	0.416667	-1.0	5.0	5.0	
2	0.562500	4.0	-6.0	-10.0	
3	0.461538	0.0	-2.0	-7.0	
4	0.500000	-2.0	1.0	5.0	

	shots_on_target_diff
0	10.0
1	5.0
2	-6.0
3	-2.0
4	2.0

[5 rows x 25 columns]

Splitting data for xgboost

```
In [49]: y = data['FTR']
```

```
x = data.drop('FTR',axis=1)
```

```
In [50]: from sklearn.model_selection import train_test_split
# Stratified split to preserve class ratios
X_train_split, X_test_split, y_train_split, y_test_split = train_test_split(
    x, y, test_size=0.2, random_state=42, stratify=y
)
```

```
In [ ]: from imblearn.over_sampling import SMOTE
from collections import Counter

from collections import Counter

# Count samples in training set
counts = Counter(y_train_split)
total = sum(counts.values())

class_weight = {cls: total / (len(counts) * count) for cls, count in counts.items()}
print(class_weight)

counts = Counter(y_train_split)
print("Before SMOTE:", counts)
undetermined_label = 1
multiplier = 2
target_samples = counts[undetermined_label] * multiplier
sm = SMOTE(sampling_strategy={undetermined_label: target_samples}, random_state=42)
X_train_res, y_train_res = sm.fit_resample(X_train_split, y_train_split)

print("After SMOTE:", Counter(y_train_res))
X_train_res = X_train_res.astype(np.float32)
y_train_res = y_train_res.astype(np.int64)
print(X_train_res.shape)
print(y_train_res.shape)

print("Class distribution after SMOTE:", dict(zip(*np.unique(y_train_res, return_counts=True))))

{1: 1.353469974829198, 2: 0.7277648878576953, 0: 1.1272836178496557}
Before SMOTE: Counter({2: 3448, 0: 2226, 1: 1854})
After SMOTE: Counter({1: 3708, 2: 3448, 0: 2226})
(9382, 24)
(9382,)
Class distribution after SMOTE: {np.int64(0): np.int64(2226), np.int64(1): np.int64(3708), np.int64(2): np.int64(3448)}
```

```
In [ ]: weights = compute_sample_weight(class_weight='balanced', y=y_train_res)
model = XGBClassifier(
    n_estimators=500,
    max_depth=5,
    learning_rate=0.05,
    objective='multi:softprob',
    num_class=3,
    eval_metric='mlogloss',
    random_state = 42
)
param_grid = {
    'n_estimators': [800, 1000],
    'max_depth': [4, 5, 6],
    'learning_rate': [0.05, 0.1],
    'subsample': [0.7, 0.8, 0.9],
    'colsample_bytree': [0.7, 0.8, 0.9]
}
grid_search = GridSearchCV(
    estimator=model,
    param_grid=param_grid,
    scoring='f1_macro',
    cv=3,
    verbose=1,
    n_jobs=-1
)

grid_search.fit(X_train_res, y_train_res, sample_weight=weights)
best_model = grid_search.best_estimator_
print("Best parameters:", grid_search.best_params_)
```

Fitting 3 folds for each of 108 candidates, totalling 324 fits
Best parameters: {'colsample_bytree': 0.9, 'learning_rate': 0.05, 'max_depth': 4, 'n_estimators': 800, 'subsample': 0.7}

```
In [ ]: from sklearn.metrics import f1_score

# Get predicted probabilities
y_proba = best_model.predict_proba(X_test_split)

# Custom thresholds per class
thresholds = [0.4, 0.6, 0.4] # start with 0.5 for all classes

def predict_with_thresholds(probs, thresholds):
    preds = []
    for p in probs:
        #Assign class if probability > threshold, else take max
        assigned = [i for i, prob in enumerate(p) if prob >= thresholds[i]]
        if assigned:
            preds.append(assigned[0])
        else:
            preds.append(np.argmax(p))
    return np.array(preds)

y_pred = predict_with_thresholds(y_proba, thresholds)

print("F1 score macro:", f1_score(y_test_split, y_pred, average='macro'))
```

F1 score macro: 0.5451007271927099

```
In [ ]: # Accuracy
print("Accuracy:", accuracy_score(y_test_split, y_pred))


# Full classification report
print(classification_report(y_test_split, y_pred))
```

Accuracy: 0.6238044633368757

	precision	recall	f1-score	support
0	0.66	0.68	0.67	556
1	0.38	0.47	0.42	464
2	0.79	0.67	0.72	862
accuracy			0.62	1882
macro avg	0.61	0.61	0.60	1882
weighted avg	0.65	0.62	0.63	1882

```
In [105]: cm = confusion_matrix(y_test_split, y_pred)

# Plot confusion matrix
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Class 0', 'Class 1', 'Class 2'],
            yticklabels=['Class 0', 'Class 1', 'Class 2'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```




```
In [106]: prec, rec, f1, _ = precision_recall_fscore_support(y_test_split, y_pred, average=None)

metrics_df = {
    "Class": ['0', '1', '2'],
    "Precision": prec,
    "Recall": rec,
    "F1-Score": f1
}

metrics_df = pd.DataFrame(metrics_df)

metrics_df.plot(x="Class", kind="bar", figsize=(8,6))
plt.title("Per-Class Metrics")
plt.ylabel("Score")
plt.xticks(rotation=0)
plt.ylim(0,1)
plt.show()
```



BALANCED RANDOM FOREST CLASSIFIER

```
In [99]: from imblearn.ensemble import BalancedRandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, f1_score
import numpy as np

# Define parameter grid
param_dist = {
    "n_estimators": [800,1000,1200],
    "max_depth": [None, 10, 20, 30],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4,5,10],
    "max_features": ['sqrt','log2']
}

# Initialize model
rf = BalancedRandomForestClassifier(random_state=42)

# Randomized search
random_search = RandomizedSearchCV(
    rf, param_distributions=param_dist,
    n_iter=20,
    cv=3,
    scoring="f1_macro",
    n_jobs=-1,
    random_state=42
)

random_search.fit(X_train_split, y_train_split)

print("Best Parameters:", random_search.best_params_)
```



```
print("Best Score (CV F1):", random_search.best_score_)
```

Best Parameters: {'n_estimators': 1000, 'min_samples_split': 10, 'min_samples_leaf': 5, 'max_features': 'sqrt', 'max_depth': 10}

Best Score (CV F1): 0.6081930932816711


```
In [100]: best_rf = random_search.best_estimator_  
  
y_proba = best_rf.predict_proba(X_test_split)  
y_pred = np.argmax(y_proba, axis=1)
```

```
In [101]: # Accuracy  
print("Accuracy:", accuracy_score(y_test_split, y_pred))  
  
# Full classification report  
print(classification_report(y_test_split, y_pred))
```

Accuracy: 0.6238044633368757


	precision	recall	f1-score	support
0	0.66	0.68	0.67	556
1	0.38	0.47	0.42	464
2	0.79	0.67	0.72	862
accuracy			0.62	1882
macro avg	0.61	0.61	0.60	1882
weighted avg	0.65	0.62	0.63	1882

```
In [102]: cm = confusion_matrix(y_test_split, y_pred)  
  
# Plot confusion matrix  
plt.figure(figsize=(6,5))  
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',  
            xticklabels=['Class 0', 'Class 1', 'Class 2'],  
            yticklabels=['Class 0', 'Class 1', 'Class 2'])  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.title("Confusion Matrix")  
plt.show()
```



```
In [103]: prec, rec, f1, _ = precision_recall_fscore_support(y_test_split, y_pred, average=None)  
  
metrics_df = {  
    "Class": ['0', '1', '2'],  
    "Precision": prec,  
    "Recall": rec,  
    "F1-Score": f1  
}  
  
metrics_df = pd.DataFrame(metrics_df)  
  
metrics_df.plot(x="Class", kind="bar", figsize=(8,6))  
plt.title("Per-Class Metrics")  
plt.ylabel("Score")  
plt.xticks(rotation=0)
```

```
plt.ylim(0,1)
plt.show()
```



RANDOM FOREST CLASSIFIER

```
In [83]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, f1_score
import numpy as np

# Define parameter grid
param_dist = {
    "n_estimators": [800,1000,1200],
    "max_depth": [None, 10, 20, 30],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4,5,10],
    "max_features": ['sqrt','log2'],
    "bootstrap": [True, False]
}

# Initialize model
rf = RandomForestClassifier(random_state=42)

# Randomized search
random_search = RandomizedSearchCV(
    rf, param_distributions=param_dist,
    n_iter=20, # number of random combinations
    cv=3, # 3-fold cross validation
    scoring="f1_macro", # you can also try "accuracy"
    n_jobs=-1,
    random_state=42
)

random_search.fit(X_train_res, y_train_res)

print("Best Parameters:", random_search.best_params_)
print("Best Score (CV F1):", random_search.best_score_)
```

```
Best Parameters: {'n_estimators': 800, 'min_samples_split': 2, 'min_samples_leaf': 2, 'max_features': 'log2', 'max_depth': None, 'bootstrap': True}
Best Score (CV F1): 0.6865774115877455
```

```
In [89]: best_rf = random_search.best_estimator_

y_proba = best_rf.predict_proba(X_test_split)
y_pred = np.argmax(y_proba, axis=1)
```

```
In [90]: # Accuracy
print("Accuracy:", accuracy_score(y_test_split, y_pred))
```


```
# Full classification report
print(classification_report(y_test_split, y_pred))
```

Accuracy: 0.6238044633368757

	precision	recall	f1-score	support
0	0.66	0.68	0.67	556
1	0.38	0.47	0.42	464
2	0.79	0.67	0.72	862
accuracy			0.62	1882
macro avg	0.61	0.61	0.60	1882
weighted avg	0.65	0.62	0.63	1882

```
In [91]: cm = confusion_matrix(y_test_split, y_pred)

# Plot confusion matrix
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Class 0', 'Class 1', 'Class 2'],
            yticklabels=['Class 0', 'Class 1', 'Class 2'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```




```
In [93]: prec, rec, f1, _ = precision_recall_fscore_support(y_test_split, y_pred, average=None)

metrics_df = {
    "Class": ['0', '1', '2'],
    "Precision": prec,
    "Recall": rec,
    "F1-Score": f1
}

metrics_df = pd.DataFrame(metrics_df)

metrics_df.plot(x="Class", kind="bar", figsize=(8,6))
plt.title("Per-Class Metrics")
plt.ylabel("Score")
plt.xticks(rotation=0)
plt.ylim(0,1)
plt.show()
```



```
In [113]: from imblearn.ensemble import EasyEnsembleClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import classification_report
import numpy as np

# Base classifier options
base_estimators = [
    DecisionTreeClassifier(max_depth=1),
    DecisionTreeClassifier(max_depth=2),
    DecisionTreeClassifier(max_depth=3)
]

# Parameter grid
param_dist = {
    "n_estimators": [10, 20, 30],
    "estimator": base_estimators
}

# Initialize EasyEnsemble
eec = EasyEnsembleClassifier(random_state=42, n_jobs=-1)

# Randomized search
rs = RandomizedSearchCV(
    estimator=eec,
    param_distributions=param_dist,
    n_iter=10,          # number of random combinations
    cv=3,              # 3-fold cross-validation
    scoring="f1_macro", # optimize macro F1 (balances all classes)
    n_jobs=-1,
    random_state=42
)

# Fit on original imbalanced data (EEC handles balancing internally)
rs.fit(X_train_split, y_train_split)

# Best parameters
print("Best Parameters:", rs.best_params_)
print("Best CV F1 (macro):", rs.best_score_)

# Evaluate on test set
y_pred = rs.predict(X_test_split)
print(classification_report(y_test_split, y_pred))
```

c:\Users\arnas\miniconda3\Lib\site-packages\sklearn\model_selection_search.py:317: UserWarning: The total space of parameters 9 is smaller than n_iter=10. Running 9 iterations. For exhaustive searches, use GridSearchCV.
warnings.warn(


Best Parameters: {'n_estimators': 10, 'estimator': DecisionTreeClassifier(max_depth=2)}

Best CV F1 (macro): 0.5960295255351604

	precision	recall	f1-score	support
0	0.69	0.60	0.64	556
1	0.35	0.56	0.43	464
2	0.79	0.60	0.68	862
accuracy			0.59	1882
macro avg	0.61	0.58	0.58	1882
weighted avg	0.65	0.59	0.61	1882

```
In [114]: cm = confusion_matrix(y_test_split, y_pred)

# Plot confusion matrix
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Class 0','Class 1','Class 2'],
            yticklabels=['Class 0','Class 1','Class 2'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
In [115]: prec, rec, f1, _ = precision_recall_fscore_support(y_test_split, y_pred, average=None)

metrics_df = {
    "Class": ['0', '1', '2'],
    "Precision": prec,
    "Recall": rec,
    "F1-Score": f1
}

metrics_df = pd.DataFrame(metrics_df)

metrics_df.plot(x="Class", kind="bar", figsize=(8,6))
plt.title("Per-Class Metrics")
plt.ylabel("Score")
plt.xticks(rotation=0)
plt.ylim(0,1)
plt.show()
```

